

Introduction to OpenCL

Concurrent Programming



OpenCL

Inhalt

- Types of parallelism
- OpenCL platform
- Programming model
 - Extensions

Types of parallelism

- Task parallelism
 - Client connection handler
 - Web crawler
 - ...
 - Data parallelism
 - Image manipulation
 - Vector arithmetics
 -
-


Types of parallelism – Task parallel hardware

- Modern CPUs
 - Multiple, full-blown CPU cores
 - Branch-prediciting
 - Caching
 - Independent from each-other (but share infrastructure)
-

Types of parallelism – Data parallel hardware

- Modern GPUs
 - SIMD – Single instruction multiple data
 - „stupid“ processors
 - One GPU processor can execute several threads at once
-

Types of parallelism – Parallel hardware



Types of parallelism – Parallel hardware

Intel® Xeon Phi™ Coprocessor Family Reference Table

| SKU # | Form Factor, Thermal | Peak Double Precision | Max # of Cores | Clock Speed (GHz) | GDDR5 Memory Speeds (GT/s) | Peak Memory BW | Memory Capacity (GB) | Total Cache (MB) | Board TDP (Watts) | Process |
|---|--------------------------------|-----------------------|---|-------------------|----------------------------|----------------|----------------------|------------------|-------------------|---------|
| SE10P <small>(special edition)</small> | PCIe Card, Passively Cooled | 1073 GF | 61 | 1.1 | 5.5 | 352 | 8 | 30.5 | 300 | 22nm |
| SE10X <small>(special edition)</small> | PCIe Card, No Thermal Solution | 1073 GF | 61 | 1.1 | 5.5 | 352 | 8 | 30.5 | 300 | |
| 5110P | PCIe Card, Passively Cooled | 1011 GF | 60 | 1.053 | 5.0 | 320 | 8 | 30 | 225 | |
| 3100 Series | PCIe Card, Actively Cooled | >1 TF | Disclosed at 3100 series launch (H1'13) | | 5.0 | 240 | 6 | 28.5 | 300 | |
| | PCIe Card, Passively Cooled | > 1 TF | | | 5.0 | 240 | 6 | 28.5 | 300 | |



PCIe Card, Actively Cooled




PCIe Card, Passively Cooled

Types of parallelism – Parallel hardware

| TECHNICAL SPECIFICATIONS | TESLA K10 ^a | TESLA K20 | TESLA K20X |
|--|--|---|----------------|
| Peak double precision floating point performance (board) | 0.19 teraflops | 1.17 teraflops | 1.31 teraflops |
| Peak single precision floating point performance (board) | 4.58 teraflops | 3.52 teraflops | 3.95 teraflops |
| Number of GPUs | 2 x GK104s | 1 x GK110 | |
| Number of CUDA cores | 2 x 1536 | 2496 | 2688 |
| Memory size per board (GDDR5) | 8 GB | 5 GB | 6 GB |
| Memory bandwidth for board (ECC off) ^b | 320 GBytes/sec | 208 GBytes/sec | 250 GBytes/sec |
| GPU computing applications | Seismic, image, signal processing, video analytics | CFD, CAE, financial computing, computational chemistry and physics, data analytics, satellite imaging, weather modeling | |
| Architecture features | SMX | SMX, Dynamic Parallelism, Hyper-Q | |
| System | Servers only | Servers and Workstations | Servers only |

Overview of work splitting

- Serial Execution



- Multi-Threading (CPU)

| | | | | |
|----|----|----|----|----|
| T0 | 0 | 1 | 2 | 3 |
| T1 | 4 | 5 | 6 | 7 |
| T2 | 8 | 9 | 10 | 11 |
| T3 | 12 | 13 | 14 | 15 |

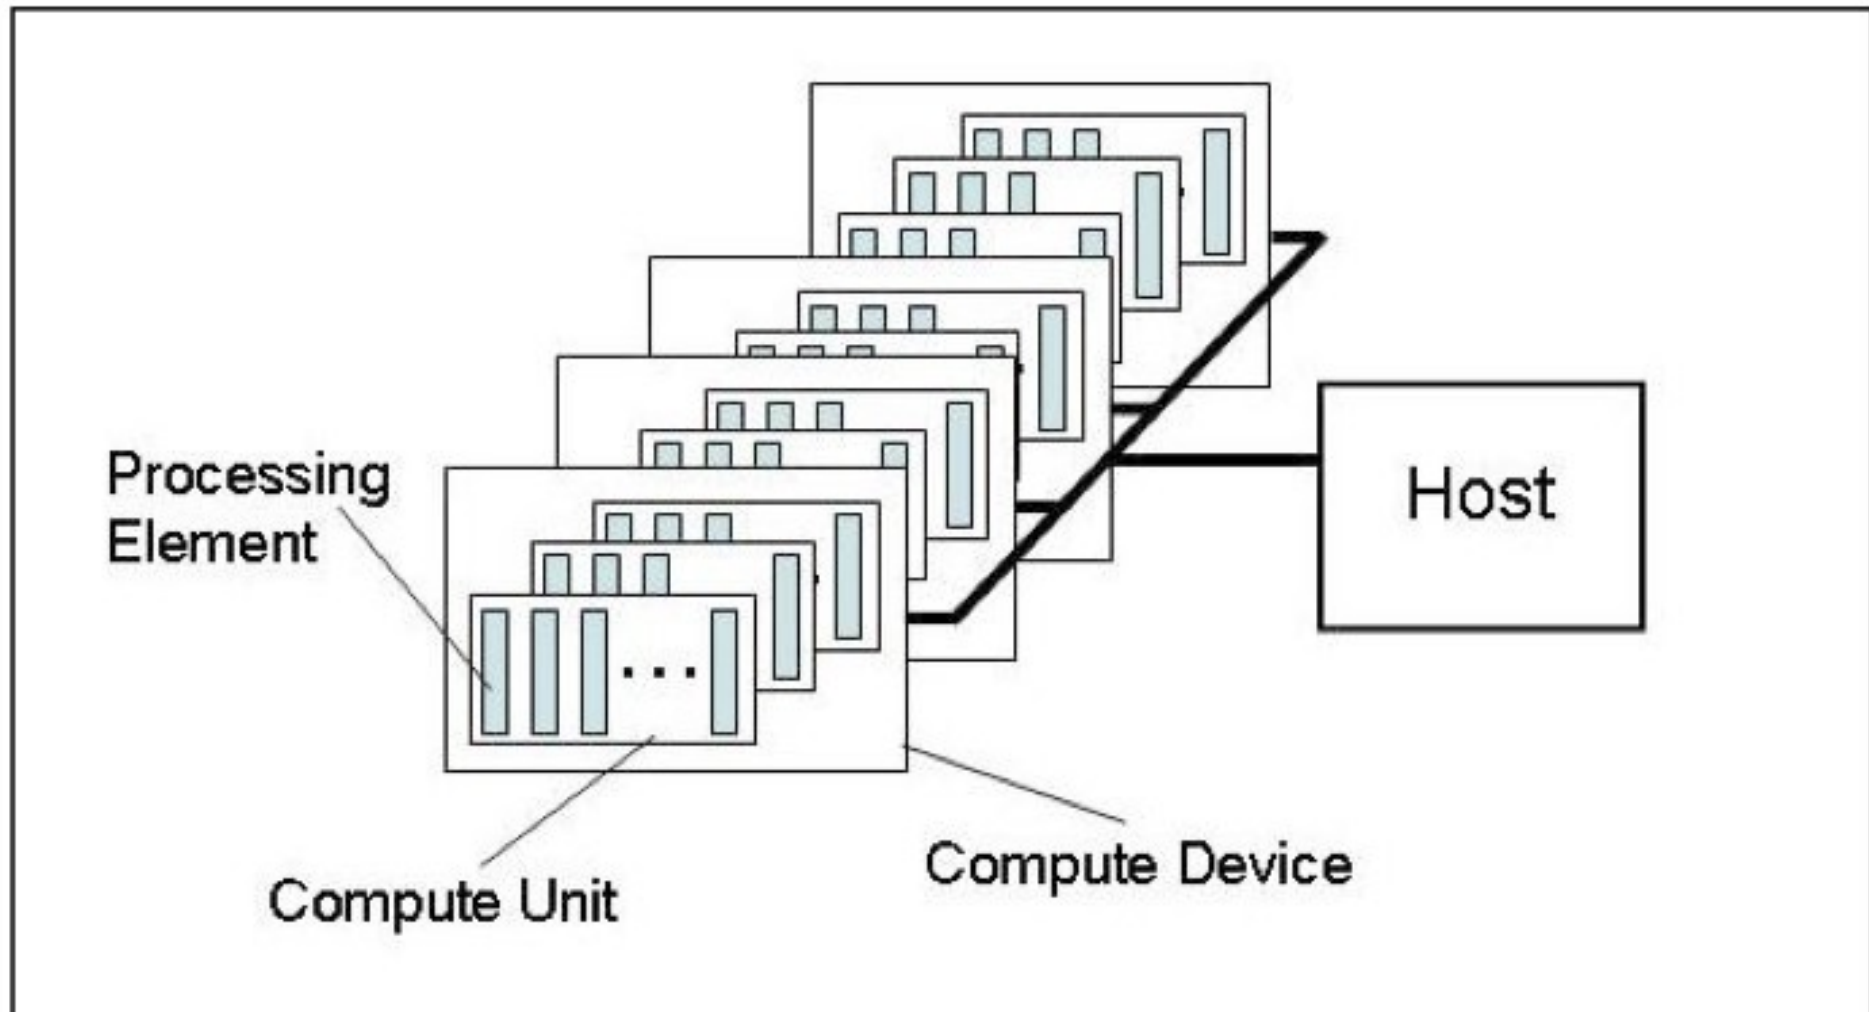
- Massivly parallel (GPU)

| | |
|-----|----|
| T0 | 0 |
| T1 | 1 |
| T2 | 2 |
| T3 | 3 |
| ... | |
| T15 | 15 |

OpenCL Platform

- Platform for data parallel applications
 - No support for task based parallelism
 - Data parallelism
 - SPMD – Single program multiple data
 - Programs are written as *Kernels*
 - One kernel executed on one Work Item
-

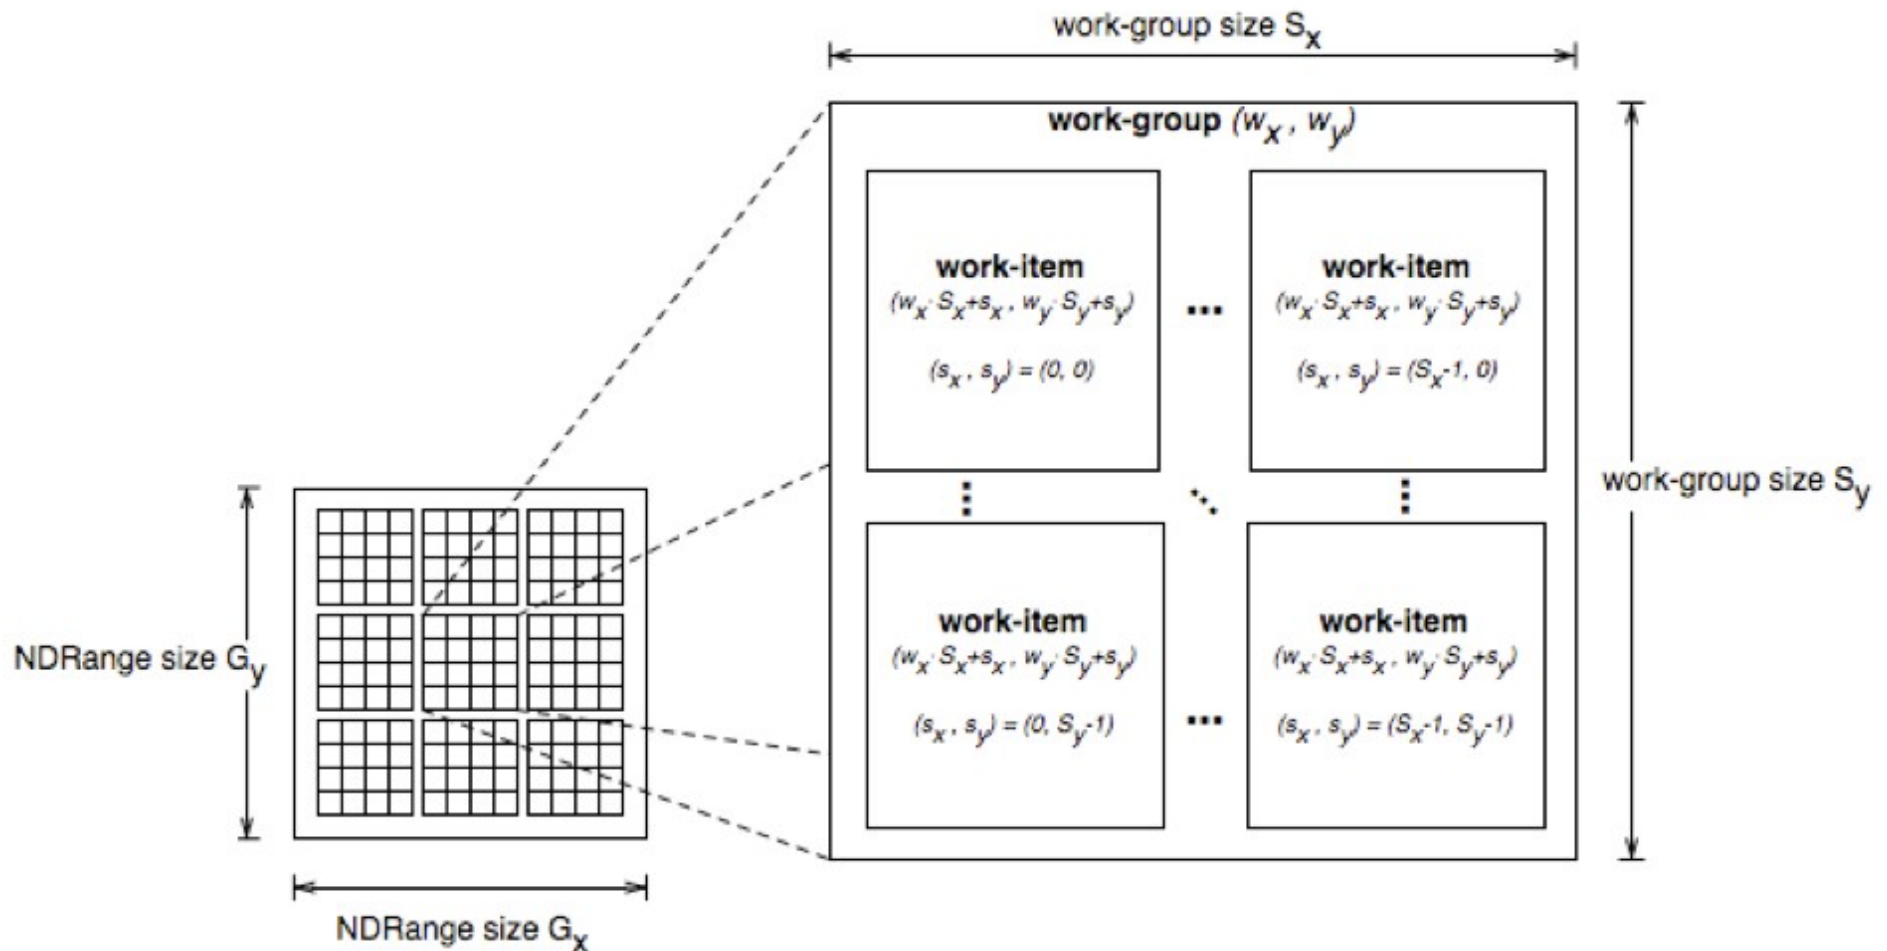
OpenCL Platform - Overview



OpenCL Platform

- Program consists of *Kernels*
 - A *Kernel* is executed on a *Work Item*
 - Work Items are grouped to *Work Groups*
-

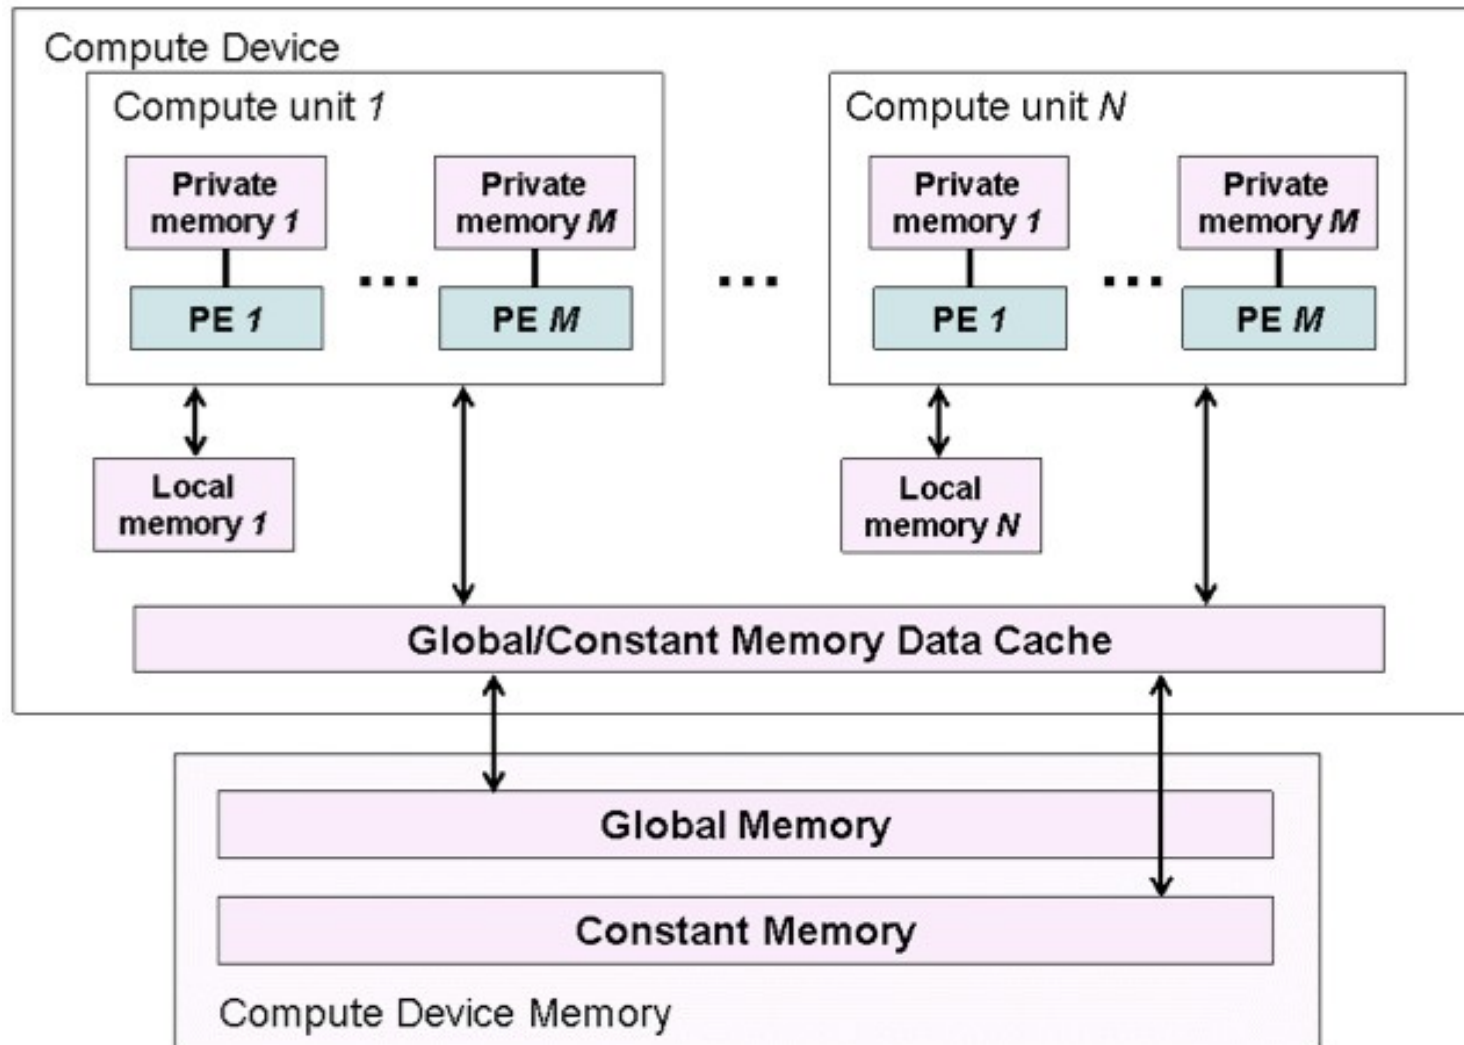
OpenCL Platform - Overview



OpenCL Platform - Memory

- Host memory
 - Memory is separated to reflect the possibilities of GPUs
 - Global memory (biggest, slowest) / Constant memory
 - All work items have access
 - On GPU: Streaming memory
 - Local memory
 - All work items in same group share it
 - Private memory (smallest, fastest)
 - Private to a work item
-

OpenCL Platform - Memory



OpenCL Platform - Memory

- Explicit memory management
 - You have to decide what to put to which memory location
- Be careful with optimizations
 - On CPUs no such thing as „local“ memory exists, only RAM

OpenCL Programming Model

- Programs are written in OpenCL C
 - Compiled at runtime by a specialized, device dependent compiler
 - Generates highly optimized code
 - Some overhead
 - No fancy memory management units on GPU
 - No SIGSEGV
 - Display server can crash
-

OpenCL Synchronization

- Barrier
 - `localBarrier()`: All kernels of a work group
 - `globalBarrier()`: Every kernel
 - Dead Lock awareness!
 - Every kernel has to have the same amount of calls to a barrier
 - Atomics
 - `cl_khr_global_int32_base_atomics`
-

Aparapi

- AMD Open Source project
 - Generates OpenCL C code from bytecode
 - Handles communication with devices
 - OpenJDK Project Sumatra
-

Aparapi – Simple Example

- Sequential version

```
final float inA[] = .... // get a float array
final float inB[] = .... // get a float array
final float result = new float[inA.length];

for (int i=0; i<array.length; i++){
    result[i]=inA[i]+inB[i];
}
```

- OpenCL

```
Kernel kernel = new Kernel(){
    @Override public void run(){
        int i= getGlobalId();
        result[i]=inA[i]+inB[i];
    }
};
Range range = Range.create(result.length);
kernel.execute(range);
```

Aparapi – OpenCL Functions

- Defined on *com.amd.aparapi.Kernel*
 - `getGlobalId(int dimension)`
 - `getLocalId(int dimension)`
 - `getGlobalSize(int dimension)`
 - `getLocalSize(int dimension)`
-

Aparapi – OpenCL Math Functions

- Defined on *com.amd.aparapi.Kernel*
 - Named like *java.lang.Math.**
 - But defined for floats as well
 - abs
 - exp
 - sqrt
 - ...
-


Aparapi – Device

- Device.best()
- Device.firstGPU()
- Device.firstCPU()
- List all OpenCL devices

```
@Test
public void printAvailableDevices() {
    OpenCLDevice.select(new DeviceSelector() {
        @Override
        public OpenCLDevice select(OpenCLDevice currentDevice) {
            System.out.println(currentDevice);
            return null;
        }
    });
}
```

Aparapi – Range

- Global Work Size has to be evenly dividable with groupSize
- `Range.create(globalWorkSize)`
- `Range.create(Device, globalWorkSize, groupSize)`
- `Range.create2D(
Device,
globalSizeD1, globalSizeD2,
groupSizeD1, groupSizeD2)`



Aparapi – Memory Management

- Annotations

```
@Local private int[] densities;  
@Constant private float[] dose;
```

- Name suffix

```
private int[] densities_$local$;  
private float[] dose_$constant$;
```

Aparapi – What's supported

- Call functions on same class
 - One dimensional arrays of primitive type
 - Some khronos extensions
 - `cl_khr_fp64`
 - `cl_khr_global_int32_base_atomics`
-

Aparapi – When something breaks

- JTP fallback

```
-----  
Jun 10, 2013 7:07:17 PM com.amd.aparapi.KernelRunner warnFallBackAndExecute  
WARNING: Reverting to Java Thread Pool (JTP) for class ch.fhnw.conpr.mandel.Ap  
com.amd.aparapi.ClassParseException: Can't assign to two dimension array
```

Aparapi

- + JVM fallback
 - + No memory management
 - + Java
 - + Debuggable
 - Not extensible
 - Code has to be rewritten
-

Links

- OpenCL Reference
 - <http://www.khronos.org/registry/cl/sdk/1.2/docs/man/xhtml/>
 - Aparapi
 - <https://code.google.com/p/aparapi/>
-

Worksheet

- Task 1
 - Mandelbrot set
 - Task 2
 - Experiment with different workloads
-